

# STA 545 Statistical Data Mining I, Fall 2020

## Homework 5

Stella Liao

October 7, 2020

1. (25 points) Please use the datasets shown in Question 1 of your homework 4 to fit a PLS model on the set A, with the parameter  $M$  chosen by the set B. Report the value of  $M$  selected by the set B, the estimated regression coefficients of the original input variables, and the test error obtained. In addition, please fit a lasso regression model on the set A, with the tuning parameter  $\lambda$  chosen by the set B. Report the test error obtained, along with the the estimated regression coefficients of the original input variables.

1) Preparing data

```
library(ISLR)
library(glmnet)
library(pls)


data(College)
set.seed(123)
train.num <- sample(777, size = 500, replace = FALSE)
train.College <- College[train.num,]
test.College <- College[-train.num,]

set.seed(1)
set.num <- sample(500, size = 250, replace = FALSE)
set.A <- College[set.num,]
set.B <- College[-set.num,]
X.set.A <- model.matrix(Apps ~ ., set.A)
X.set.B <- model.matrix(Apps ~ ., set.B)
test.matrix.College <- model.matrix(Apps ~ ., test.College)
grid <- 10 ^ seq(10, -2, length = 100)
```

2) PLS model

```
fit.PLS <- plsr(Apps ~ .,
               data = set.B,
               family = "gaussian",
               scale = F,
               validation = "CV")
validationplot(fit.PLS, val.type = "MSEP")
```

## Apps



```
summary(fit.PLS)
```

```
## Data:      X dimension: 527 17
## Y dimension: 527 1
## Fit method: kernelpLS
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           4156    2038    1983    1675    1480    1404    1406
## adjCV         4156    2033    1981    1668    1467    1396    1397
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV           1406    1396    1396    1306    1287    1271    1263
## adjCV         1397    1388    1388    1298    1279    1260    1254
##      14 comps 15 comps 16 comps 17 comps
## CV           1259    1260    1260    1261
## adjCV         1250    1251    1251    1251
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X           40.42    68.55    92.44    97.05    98.76    99.29    99.65    99.97
## Apps        77.46    82.01    86.86    90.42    90.81    90.86    90.93    90.97
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps 15 comps
## X           100.00    100.00    100.00    100.00    100.0    100.00    100.00
## Apps        91.12    92.18    92.49    92.89    92.9    92.91    92.92
##      16 comps 17 comps
## X           100.00    100
## Apps        92.92    93
```

According to the output of `summary(fit.pcr2)`, we find that when `ncomp = 16`, the cross validation error is lowest.

```
fit.PLS2 <- pcr(Apps ~ .,
               data = set.A,
               family = "gaussian",
               scale = F,
               ncomp = 16)

pred.PLS <- predict(fit.PLS2,
                   test.College,
                   scale = F,
                   ncomp = 16)

# test error in PLS model
test.error.PLS <- mean((pred.PLS - test.College$Apps)^2)
test.error.PLS
```

```
## [1] 1726332
```

```
#the coefficients of the original outputs in the PLS model
as.data.frame(fit.PLS2$coefficients[, , 16])
```

```
##               fit.PLS2$coefficients[, , 16]
## PrivateYes           -1.16873440
## Accept                1.31283568
## Enroll               -0.44554162
## Top10perc            13.32145875
## Top25perc             2.71971884
## F.Undergrad           0.05365231
## P.Undergrad           0.09952674
## Outstate             -0.03998232
## Room.Board            0.13036972
## Books                 -0.23101502
## Personal              -0.10137072
## PhD                   -2.75395703
## Terminal              -4.59517609
## S.F.Ratio             39.73244558
## perc.alumni           -7.99349032
## Expend                 0.09091874
## Grad.Rate              1.56692304
```

3) Lasso regression model

```
lasso.mod <- glmnet(X.set.A,
                   set.A$Apps,
                   family = "gaussian",
                   standardize = F,
                   alpha = 1,
                   lambda = grid)

cv.out = cv.glmnet(X.set.B,
                  set.B$Apps,
                  lambda = grid,
                  standardize = F,
                  alpha = 1)
```

```

bestlam.B.lasso = cv.out$lambda.min
bestlam.B.lasso

## [1] 231.013

lasso.pred = predict(lasso.mod,
                      s = bestlam.B.lasso,
                      newx = test.matrix.College)

# the test error in the lasso model
mean((lasso.pred - test.College$Apps)^2)

## [1] 1777035

#the coefficients of the original outputs in the lasso model
lasso.coef = predict(lasso.mod,
                     type = "coefficients",
                     s = bestlam.B.lasso)

lasso.coef

## 19 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) -516.12632721
## (Intercept) .
## PrivateYes .
## Accept      1.29229257
## Enroll      -0.33117710
## Top10perc   11.48847212
## Top25perc   2.38685707
## F.Undergrad 0.04720434
## P.Undergrad 0.10098330
## Outstate    -0.04880593
## Room.Board  0.14189798
## Books       -0.20688684
## Personal    -0.11410355
## PhD         -1.15822172
## Terminal    -2.93639397
## S.F.Ratio   8.28915054
## perc.alumni -6.30300183
## Expend      0.08161484
## Grad.Rate   0.10310942

```

## Problem 2


Define the loss function as.

$$f(\beta_j) = (\beta_j - \hat{\beta}_j^{ols})^2 + \lambda \alpha \beta_j^2 + \lambda(1-\alpha) |\beta_j|$$

1) when  $\beta_j \geq 0$ .

$$\begin{aligned} f(\beta_j) &= (\beta_j - \hat{\beta}_j^{ols})^2 + \lambda \alpha \beta_j^2 + \lambda(1-\alpha) \beta_j \\ &= (\lambda \alpha + 1) \beta_j^2 + [\lambda(1-\alpha) - 2\hat{\beta}_j^{ols}] \beta_j + (\hat{\beta}_j^{ols})^2 \end{aligned}$$


the graph of  $f(\beta_j)$  would be like and  $\hat{\beta}_j^{ols} > \frac{\lambda(1-\alpha)}{2}$



2) when  $\beta_j < 0$ .

$$\begin{aligned} f(\beta_j) &= (\beta_j - \hat{\beta}_j^{ols})^2 + \lambda \alpha \beta_j^2 - \lambda(1-\alpha) \beta_j \\ &= (\lambda \alpha + 1) \beta_j^2 - [\lambda(1-\alpha) + 2\hat{\beta}_j^{ols}] \beta_j + (\hat{\beta}_j^{ols})^2 \end{aligned}$$

the graph of  $f(\beta_j)$  would be like and  $\hat{\beta}_j^{ols} < \frac{-\lambda(1-\alpha)}{2}$



Therefore, the solution is

$$\hat{\beta}_j = \begin{cases} \frac{2\hat{\beta}_j^{ols} - \lambda(1-\alpha)}{2(\lambda \alpha + 1)} & \text{if } \hat{\beta}_j^{ols} > \frac{\lambda(1-\alpha)}{2} \\ 0 & \text{if } \hat{\beta}_j^{ols} \leq \left| \frac{\lambda(1-\alpha)}{2} \right| \\ \frac{\lambda(1-\alpha) + 2\hat{\beta}_j^{ols}}{2(\lambda \alpha + 1)} & \text{if } \hat{\beta}_j^{ols} < \frac{-\lambda(1-\alpha)}{2} \end{cases}$$

And we could also use the coordinate descent algorithm for the lasso method to solve it.

**Problem 3 (50 points)** Please write your own R function for the coordinate descent algorithm to fit lasso regression models. Please use the prostate cancer data to compare the results from your own R function with the results from the glmnet R function in the glmnet R package.

1) using glmnet() function

```
prostate <- read.csv("prostate.csv")
# have a look at prostate data
head(prostate, n = 2 )

##      X      lcavol  lweight age      lbph svi      lcp gleason pgg45      lpsa
## 1 1 -0.5798185 2.769459 50 -1.386294 0 -1.386294      6      0 -0.4307829
## 2 2 -0.9942523 3.319626 58 -1.386294 0 -1.386294      6      0 -0.1625189
##      train
## 1      TRUE
## 2      TRUE

# standardizing prostate data to make sure its mean equal to 0 and variance equal to 1
X.prostate = scale(as.matrix(prostate[,2:9]))
y.prostate = scale(as.vector(prostate[,10]))

set.seed(123)
lasso.mod.prostate <- glmnet(X.prostate,
                             y.prostate,
                             intercept = F,
                             standardize = T,
                             family="gaussian",
                             alpha = 1,
                             lambda = grid)

cv.out.prostate = cv.glmnet(X.prostate,
                             y.prostate,
                             alpha = 1,
                             standardize = T,
                             intercept = F,
                             lambda = grid)

bestlam.B.lasso.prostate = cv.out.prostate$lambda.min
coef.lasso.prostate = predict(lasso.mod.prostate, type = "coefficients", s = bestlam.B.lasso.prostate)
coef.lasso.prostate

## 9 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  .
## lcavol      0.54796374
## lweight     0.22164654
## age        -0.10768686
## lbph        0.10690468
## svi         0.24463535
## lcp        -0.06094245
## gleason     0.02015517
## pgg45       0.08356488
```

```
bestlam.B.lasso.prostate
```

```
## [1] 0.01
```

2) using coordinate descent algorithm [1]

The loss function for lasso problem is

$$f(\beta) = \frac{1}{2} \sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

According to the coordinate algorithm,

$$\min_{\beta_k} f(\beta_1, \beta_2, \dots, \beta_k, \beta_{k+1}, \dots, \beta_p)$$

$$\begin{aligned} \frac{\partial f(\beta)}{\partial \beta_k} &= - \sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij} \beta_j) x_{ik} + \frac{\partial (\lambda |\beta_k|)}{\partial \beta_k} \\ &= - \sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij} \beta_j) x_{ik} + \sum_{i=1}^n x_{ik}^2 \beta_k + \frac{\partial (\lambda |\beta_k|)}{\partial \beta_k} \\ &= -r_k + z_k \beta_k + \frac{\partial (\lambda |\beta_k|)}{\partial \beta_k} \end{aligned}$$

$$r_k = \sum_{i=1}^n (y_i - \sum_{j=1}^p x_{ij} \beta_j) x_{ik}$$

$$z_k = \sum_{i=1}^n x_{ik}^2$$

And easily, we could know that

$$\frac{\partial (\lambda |\beta_k|)}{\partial \beta_k} = \begin{cases} \lambda & , \beta_k > 0 \\ [-\lambda, \lambda] & , \beta_k = 0 \\ -\lambda & , \beta_k < 0 \end{cases} \Rightarrow \frac{\partial f(\beta)}{\partial \beta_k} = \begin{cases} -r_k + z_k \beta_k + \lambda, & \beta_k > 0 \\ [-r_k - \lambda, -r_k + \lambda], & \beta_k = 0 \\ -r_k + z_k \beta_k - \lambda, & \beta_k < 0 \end{cases}$$

we make  $\frac{\partial f(\beta)}{\partial \beta_k} = 0$ , then we could get the minimum value

$$\beta_k^* = \begin{cases} (r_k - \lambda) / z_k & , r_k > \lambda \\ 0 & , r_k < |\lambda| \\ (r_k + \lambda) / z_k & , r_k < -\lambda \end{cases}$$

In the soft-thresholding function  $S(a, b) = \begin{cases} a - b, & a > b \\ 0, & a < |b| \\ a + b, & a < -b \end{cases}$

$$\beta_k^* = \frac{1}{z_k} S(r_k, \lambda)$$



```

cd.lasso <- function(X, y, lambda = 0.1, max.iter = 1000, tol = 1e-6 ){

  #to create soft-thresholding function
  soft.thresholding <- function(b, l){
    result = rep(0, length(b))
    result[b > l] = b[b > l] - l
    result[b < -l] = b[b < -l] + l
    result
  }

  #to create function to calculate beta k star
  compute.bks <- function(k, X, y, beta, lambda){
    y.predict = X %*% beta
    rk = X[, k] %*% (y - y.predict + X[, k] * beta[k])
    rk = rk / nrow(X)
    zk = colSums(X^2)[k]
    zk = zk / nrow(X)
    beta.k = soft.thresholding(rk, lambda)
    beta.k = beta.k / zk
    beta.k }

  #initialize some parameters
  tol.curr = 1
  iter = 1
  all.beta = rep(0, ncol(X))
  old.all.beta = rep(0, ncol(X))

  #update beta k
  while (tol < tol.curr && iter < max.iter) {
    for (k in 1:ncol(X)) {
      old.all.beta[k] = all.beta[k]
      all.beta[k] = compute.bks(k, X, y, all.beta, lambda)
    }
    tol.curr = abs(all.beta - old.all.beta)
    iter = iter + 1
  }
  #print all beta
  all.beta
}

coef.cd.lasso.prostate <- cd.lasso(X = X.prostate,
                                   y = y.prostate,
                                   lambda = bestlam.B.lasso.prostate,
                                   tol = 1e-12)

comparison <- data.frame("glmnet" = as.data.frame(summary(coef.lasso.prostate))$x,
                         "coordinate descent" = coef.cd.lasso.prostate)

predictors <- c("lcavol", "lweight", "age", "lbph", "svi", "lcp", "gleason", "pgg45")
comparison <- cbind(predictors, comparison)

comparison

## predictors      glmnet coordinate.descent
## 1      lcavol  0.54796374      0.54775369

```

## 2	lweight	0.22164654	0.22161781
## 3	age	-0.10768686	-0.10755924
## 4	lbph	0.10690468	0.10682744
## 5	svi	0.24463535	0.24451606
## 6	lcp	-0.06094245	-0.06062363
## 7	gleason	0.02015517	0.02016581
## 8	pgg45	0.08356488	0.08342290

We could see that the coefficients of those predictors in the two functions are almost same.

## Reference

[1] [https://www.scutmath.com/coordiante\\_descent\\_for\\_lasso.html](https://www.scutmath.com/coordiante_descent_for_lasso.html)